

# Prácticas BigData

## 1. Administración de HDFS

- Realizar un report del sistema HDFS

### **hdfs dfsadmin -report**

Configured Capacity: 50940104704 (47.44 GB)  
 Present Capacity: 44860829696 (41.78 GB)  
 DFS Remaining: 44859387904 (41.78 GB)  
 DFS Used: 1441792 (1.38 MB)  
 DFS Used%: 0.00%  
 Under replicated blocks: 0  
 Blocks with corrupt replicas: 0  
 Missing blocks: 0

### ----- Live datanodes (1):

Name: 127.0.0.1:50010 (localhost)  
 Hostname: localhost  
 Decommission Status : Normal  
 Configured Capacity: 50940104704 (47.44 GB)  
 DFS Used: 1441792 (1.38 MB)  
 Non DFS Used: 6079275008 (5.66 GB)  
 DFS Remaining: 44859387904 (41.78 GB)  
 DFS Used%: 0.00%  
 DFS Remaining%: 88.06%  
 Configured Cache Capacity: 0 (0 B)  
 Cache Used: 0 (0 B)  
 Cache Remaining: 0 (0 B)  
 Cache Used%: 100.00%  
 Cache Remaining%: 0.00%  
 Xceivers: 1  
 Last contact: Mon Apr 20 08:17:15 CEST 2015

- Comprobar con hdfs fsck el estado del sistema de ficheros

### **hdfs fsck /**

```

Connecting to namenode via http://localhost:50070
FSCK started by hadoop (auth:SIMPLE) from /127.0.0.1 for path / at Mon Apr 20
08:19:57 CEST 2015
.....Status: HEALTHY
Total size: 1380389 B
Total dirs: 6
Total files: 6
Total symlinks: 0
Total blocks (validated): 5 (avg. block size 276077 B)
Minimally replicated blocks: 5 (100.0 %)
Over-replicated blocks: 0 (0.0 %)
Under-replicated blocks: 0 (0.0 %)
Mis-replicated blocks: 0 (0.0 %)
Default replication factor: 1
Average block replication: 1.0
Corrupt blocks: 0
Missing replicas: 0 (0.0 %)
Number of data-nodes: 1
Number of racks: 1
FSCK ended at Mon Apr 20 08:19:57 CEST 2015 in 14 milliseconds

The filesystem under path '/' is HEALTHY

```

- También Podemos comprobar el estado de un determinado directorio, por ejemplo de prácticas

```

hdfs fsck /practicass
Connecting to namenode via
http://localhost:50070/fsck?ugi=hadoop&path=%2Fpracticass
FSCK started by hadoop (auth:SIMPLE) from /127.0.0.1 for path /practicass at Sat
Jan 06 19:55:19 CET 2018
...Status: HEALTHY
Total size: 2647821 B
Total dirs: 2
Total files: 3
Total symlinks: 0
Total blocks (validated): 2 (avg. block size 1323910 B)
Minimally replicated blocks: 2 (100.0 %)
Over-replicated blocks: 0 (0.0 %)

```

```
Under-replicated blocks: 0 (0.0 %)
Mis-replicated blocks: 0 (0.0 %)
Default replication factor: 1
Average block replication: 1.0
Corrupt blocks: 0
Missing replicas: 0 (0.0 %)
Number of data-nodes: 1
Number of racks: 1
FSCK ended at Sat Jan 06 19:55:19 CET 2018 in 7 milliseconds
```

- Podemos ver la topología que tenemos en este momento

```
hdfs dfsadmin -printTopology
```

```
Rack: /default-rack
```

```
127.0.0.1:50010 (localhost)
```

- Podemos comprobar si hay algún fichero abierto

```
hdfs dfsadmin -listOpenFiles
```

Client Host	Client Name	Open File Path
-------------	-------------	----------------

-